

The PDF extraction pipeline in gieok normalises text on a per-page basis before handing it to the shared chunker, repairing line-wrapping artefacts that a justified PDF page introduces, such as words broken across a hyphenated line ending, so retrieval quality does not degrade for text extracted from a PDF.

Without paragraph reconstruction a wrapped PDF page arrives at the chunker as one single giant paragraph and falls straight through to the hard character-window split, which is the lowest-quality retrieval boundary the chunker has to offer.